



Chapter 1

TCP

Congestion Control

Outline


- Connection-oriented Transport: TCP
- Segment Structure
- Connection Management
- Reliable Data Transfer

TCP Overview

- Point-to-point
 - One sender, one receiver
- Reliable
 - Segments delivered in-order without loss
- In-order byte stream
 - No message / record boundaries
- Pipelined
 - Sliding-window type control algorithm

TCP Overview

- Full duplex data
 - bi-directional data flow in the same connection
- Connection-oriented
 - handshaking (exchange of control messages) initializes sender & receiver state before data exchange



- Both sides have buffers
 - Lots of “Producer-Consumer” coordination problems to overcome

Transmission Control


- Flow Control
 - Sender will not overwhelm receiver
- Congestion Control
 - Various algorithms employed to limit sending of segments
 - Don't want to overwhelm the network

Outline

- Connection-oriented Transport: TCP
- Segment Structure
- Connection Management
- Reliable Data Transfer


Format

- Source port #: Identifies the **sending application** (process) on the source host.
- Example: HTTP client might use port 49152.
- Destination port: Identifies the **receiving application** on the destination host.
- Example: HTTP server listens on port 80 (or 443 for HTTPS).




Format

- Seq# and Ack#
 - counted by bytes of data, not segments
- Seq: number of the first byte in this segment
- Ack: number of the NEXT expected byte
 - Acks are cumulative
 - Example: ACK = 101 means bytes up to 100 have been received correctly.




Flag Field

- Header Length (4 bits)
 - # words (32-bit)
 - Minimum values: 5 (20 bytes)
 - Can be longer with options




Flag Field

- ACK: Segment acknowledges receipt of another segment (i.e. ack seq# is good)
- Setup and teardown signals
 - SYN: Synchronize seq#
 - FIN: No more data from sender
 - RST: Reset connection
- Rarely used
 - PSH: Push the data to app-layer immediately
 - URG: Urgent data, indicated by Urgent Data Pointer




Format

- Receive window
 - #bytes receiver is willing to accept
 - For flow-control




Format

- Checksum
 - Error-checking for header + data.
 - Ensures integrity.
- Urgent data pointer
 - Offset in data field
- Options
 - Time stamps
 - Window scaling factors
 - Negotiating MSS



Format

- Application Data
 - Size limited by MSS
- MSS=Maximum Segment Size
 - Despite name, MSS is most app data that can be carried
- MTU (Max Trans Unit) of lower level generally drives MSS
 - App data + TCP header + IP overhead must fit in MTU
 - MSS is often 1460, 536 or 512 bytes



Outline


- Connection-oriented Transport: TCP
- Segment Structure
- Connection Management
- Reliable Data Transfer

TCP Connection Management

- Why connection establishment?
 - TCP sender, receiver setup state before exchanging data segments
 - Initialize TCP variables:
 - ❖ seq#s
 - ❖ buffers, flow control info (*RcvWindow*)
- Why requiring connection teardown?
 - **Free up state** – remove connection info from memory (control blocks, buffers, sequence tracking).
 - **Release resources** – so ports, memory, and buffers can be reused for new connections.
 - **Ensure reliability** – confirm both sides agree that the connection is finished (no data left in transit).

Three-way Handshake

- Step 1: client sends TCP SYN segment to server
 - specifies initial seq#
 - holds no data
- Step 2: server responds with SYNACK segment
 - server allocates buffers
 - specifies initial seq#
- Step 3: Client replies with ACK
 - May contain data



Question

- TCP specification requires each side of a connection to select an initial starting sequence number at random¹. Why?
 - Avoid confusion with old connections
 - Security benefit
 - ❖ Predictable ISNs (like always starting at 0 or incrementing by 1) can be exploited for TCP session hijacking.
 - ❖ Random ISNs make it harder for attackers to inject fake packets into a connection.

Question

- Let's say TCP does not exchange initial sequence numbers, and just use 0 as the starting point. What can happen?
 - A client connects to a server → sends data (seq numbers 0,1,2,...).
 - Connection closes, but some stray packets are still in the network (delayed).
 - Client reconnects immediately to the same server (same IP/port tuple).
 - New connection also starts at 0 → delayed old packets fall into the valid sequence number range and get accepted as new data.

Result:


- Data corruption or wrong application behavior.

Example

- If a host receives a TCP SYN segment to a closed port (no process is listening)
 - The host replies with a TCP RST (Reset) segment.
 - This tells the sender:
“That port is closed — no connection will be established.”
 - RST is TCP’s way of signaling an error condition.


Closing a Connection

- Step 1: Client sends TCP FIN segment to server
- Step 2: Server receives FIN, responds with ACK
 - Closes connection
 - Sends FIN
- Recall: Connection is bi-directional, needs to be shut down from each side




Closing a Connection

- Step 3: client receives FIN, replies with ACK
 - Client enters “timed wait”
 - Will respond to FIN with ACK
 - 240 seconds ($2 * \text{max segment lifetime} - \text{MSL}^1$)
- Step 4: server receives ACK, closes connection
- ¹MSL = longest time a TCP segment is allowed to exist in the network before being discarded.
- Typical MSL = 120 seconds



Question


- Why does client enter “timed wait” state before “closed,” even after receiving FIN from server?



Answer

Result:

- If ACK from the client is lost, and the client does not wait?
 - Client may open the same connection again (same pair of port#s)
 - Receives FIN from earlier incarnation of connection
 - Immediately initiate closing of the later incarnation
- Any stray packets from the old connection have expired ($1 \times MSL$).



Outline


- Connection-oriented Transport: TCP
- Segment Structure
- Connection Management
- Reliable Data Transfer (RDT)

TCP RDT

- TCP creates RDT service on top of IP's unreliable service
 - Pipelined segments
 - Cumulative acks
 - Retransmission timer
- Retransmissions are triggered by:
 - Timeout events
 - Duplicate acks
- Initially, we consider simplified TCP sender:
 - Ignore duplicate acks
 - Ignore flow control, congestion control
 - assume RTT is estimated somehow

TCP Seq# and ACKs


- Sequence numbers:
 - byte stream “number” of first byte in segment’s data
- ACKs:
 - seq # of next byte expected from other side
 - cumulative ACK
 - ❖ acknowledges bytes up to the first missing byte in the stream
 - piggybacked
 - ❖ with data, if possible



TCP Sender Events

- Data received from app:
 - Create segment
 - seq# is byte-stream number of first data byte in segment
 - Send, if allowed by congestion & flow-control
 - start timer if not already running (**think of timer as for oldest unacked segment**)
 - expiration interval:
TimeOutInterval
- Timeout:
 - retransmit segment that caused timeout
 - restart timer
- ACK received:
 - If acknowledges previously unACKed segments
 - update what is known to be ACKed
 - start timer if there are outstanding segments


Retransmission Scenarios



Lost ACK

- What does server do when it gets the retransmitted segment?
- Discard it! Expected seq# is 100. Since $92 < 100$ it knows this is duplicate data


Retransmission Scenarios



Premature Timeout

- Will client retransmit 2nd segment ($\text{Seq}=100$)?
- No! $\text{ACK}=120$ arrives before the timeout

Retransmission Scenarios



Cumulative ACK

- Host hasn't received ACK for first segment, so why doesn't it retransmit?
- ACK=120 is cumulative, which means everything up to byte 120 has been received (including 92)


ACK Generation

Event at receiver	Receiver Action
Arrival of in-order segment with expected seq#. All data up to seq# already ACKed	<i>Delayed ACK.</i> Wait up to 500ms for next segment. If no next segment, send ACK
Arrival of in-order segment with expected seq#. One other segment has ACK pending	Immediately send single cumulative ACK for both in-order segments
Arrival of out-of-order segment higher than expected seq#. Gap detected	Immediately send <i>duplicate ACK</i> , indicating seq# of next expected byte
Arrival of segment that partially or completely fills gap	Immediately send ACK, provided that segment starts at lower end of gap

Fast Retransmit

- Time-out period often relatively long:
 - Long delay before resending lost segment
- Detect lost segments via duplicate ACKs
 - Sender often sends many segments back-to-back
 - If segment is lost, there will likely be many duplicate ACKs
- If sender receives 3 duplicate ACKs, it supposes that segment after ACKed data was lost:
- **Fast retransmit:** resend segment before timer expires

Scenario



Fast Retransmission

- Sender doesn't have to wait for a timeout to notice probable loss of seq=100 segment
- Sort of a NACK

Why 3?

- Why 3 duplicate ACKs? Why not do fast retransmit after the first duplicate ACK for a segment is received?
- If $n+1$ and $n+2$ (or $n+3$) are just reordered, then waiting for 2 duplicate ACKs will not retrigger retransmission
 - Voodoo constant

